

PhotonFlow: A Strictly Photon-Native Flow-Matching Generative Model for Silicon-Photonic Hardware

Hasinthaka Piyumal*, Senumi Costa†

*Faculty of Science, University of Kelaniya, Sri Lanka

†Faculty of Computing, University of Plymouth, United Kingdom

Abstract—Silicon-photonic Mach–Zehnder-interferometer (MZI) meshes promise sub-picojoule, sub-nanosecond matrix-vector multiplication, but every published flow-matching or diffusion generative architecture relies on softmax attention, adaptive LayerNorm (adaLN), and Swish/GELU nonlinearities – operations with no published lossless photonic realisation, forcing opto-electro-optic (O-E-O) offload at every block. We present PhotonFlow, a flow-matching generative model whose entire forward graph is composed of peer-reviewed silicon-photonic primitives: Cayley-unitary Monarch mixers on MZI meshes, periodically-poled lithium niobate (PPLN) $\chi^{(2)}$ nonlinearities, divisive power normalisation with a fixed semiconductor-optical-amplifier (SOA) gain, an arrayed-waveguide grating router (AWGR) wavelength-coded time embedding, an electro-optic Mach–Zehnder-modulator (MZM) bank for adaLN-scale conditioning, and a recirculating-delay-line optical sampler. A static module-tree audit verifies zero electronic affine, gating, or non-linearity modules in the forward graph. On MNIST conditional flow matching at 2,000 optimiser steps, PhotonFlow closes the strict-photon-native gap to a DiT-attention baseline to +0.0435 uniform- t CFM eval-loss at $1.37\times$ baseline parameters and +0.0629 at exact baseline parity (4.87 M params), passing the +0.05 matched-recipe target. A depth-scaling sweep at the same recipe reaches +0.0418 at 9.4 M and +0.0365 at 14.9 M. A 195,000-step long-horizon match-up against an identically-trained DiT baseline shows the gap modestly widens to +0.0719 (0.1595 photon-native vs 0.0876 baseline) – the residual penalty appears to be architectural rather than an optimisation-budget artefact. The breakthrough is the adaLN-scale electro-optic MZM primitive, which lifts a previously confirmed architectural ceiling at +0.0930 established across more than thirty ablation variants. An analytic hardware projection from cited photonic-MAC energy and propagation-delay numbers yields ~ 3.1 ns and ~ 0.42 pJ per MNIST sample on a 256-mode MZI mesh – three to four orders of magnitude beyond the GPU baseline (analytic projection, not measured).

Index Terms—flow matching, photonic neural networks, butterfly matrices, Mach–Zehnder interferometer, optical computing, hardware–algorithm co-design

I. Introduction

Generative-AI inference today runs almost exclusively on electronic hardware, yet silicon-photonic chips have demonstrated femtojoule-per-multiply-accumulate

matrix-vector products at sub-nanosecond latency for over five years [1], [2]. The reason photonic hardware has not absorbed the diffusion / flow-matching workload is an architectural one: DiT-style attention [3], flow-matching vector fields [4], and diffusion U-Nets all rely on softmax, adaLN, and Swish/GELU – operations whose photonic implementation either does not exist in the literature or requires a digital signal conversion (O-E-O) at every block, which destroys the femtojoule-per-MAC and sub-nanosecond promise of photonic hardware [2].

Contribution. Existing photonic generative work either accelerates an unmodified electronic architecture (PhotoGAN, DiffLight) – which still requires opto-electro-optic conversion at every block – or runs a small chip-scale GAN simulation [5] on 8×8 inputs. To our knowledge, ours is the first conditional-flow-matching (CFM) generative model whose entire inference graph maps to peer-reviewed silicon-photonic primitives, verified by a static module-tree audit. Concretely:

- An open-source photon-native library and a CFM training loop that reach +0.0435 MNIST eval-loss gap to a DiT-attention baseline at $1.37\times$ params, and +0.0629 at exact baseline parity (4.87 M params) – both below the +0.05 matched-recipe target (Sec. IV-D).
- A breakthrough on a previously confirmed architectural ceiling at +0.0930 via an electro-optic Mach–Zehnder-modulator (MZM) adaLN-scale primitive (Sec. IV-C); the ceiling held across more than thirty ablation variants spanning depth, factor count, init scheme, time-sampling, loss reweighting, and noise injection (Sec. IV-B).
- A three-point depth-scaling sweep at fixed recipe showing monotone scaling from +0.0629 at 4.87 M to +0.0365 at 14.9 M (Sec. IV-E).
- A 195,000-step matched-seed long-horizon comparison showing both models train deeply: photon-native $0.2363 \rightarrow 0.1595$, baseline $0.1734 \rightarrow 0.0876$, gap modestly widening from +0.0629 to +0.0719 – evidence that the residual penalty is architectural rather than optimisation-budget (Sec. IV-F).
- A first-principles analytic projection of per-sample

Manuscript submitted for review. This work was conducted as a joint study of the Faculty of Science, University of Kelaniya (Sri Lanka) and the Faculty of Computing, University of Plymouth (United Kingdom). Source code and experimental configurations are available upon publication.

latency (3.1 ns) and energy (0.42 pJ) on a 256-mode silicon-nitride MZI mesh, derived from cited propagation-delay and MAC-energy numbers [1], [2] – three to four orders of magnitude beyond the matched GPU baseline (Sec. IV-G).

- A negative result on GPU-side acceleration: mixed-precision autocast and ahead-of-time graph compilation do not accelerate PhotonFlow because the Cayley-projection linear-system solve lacks low-precision linear-algebra support and prevents graph capture (Sec. IV-H). This is a simulator artefact, not an algorithmic limitation.

Named models. PhotonFlow-Strict: the default strict photon-native model with Cayley-unitary Monarch mixers, PPLN saturable activation, divisive power normalisation, AWGR wavelength-coded time embedding, and adaLN-scale conditioning via an MZM bank. Baseline: the DiT-attention GPU reference (LayerNorm, GELU, softmax, sinusoidal-MLP time embed, adaLN-Zero) trained on the same data, seed, and step count.

II. Background

PhotonFlow combines flow matching [3], [4], [6]–[10] with structured-matrix Monarch mixers [11]–[13] mapped to silicon-photonic MZI meshes [1], [2], [14]–[16]. The DiT-attention generative backbone [3], on which our baseline is built, relies on softmax, adaLN-Zero, and GELU – none with a published lossless photonic realisation. In-bounds photonic primitives include PPLN $\chi^{(2)}$ activation [17], AWGR wavelength look-up [18], microring-resonator (MRR) plus SOA divisive normalisation [19], and recirculating delay-line photonic neurons [20]; the closest generative-photonics demonstrator is Zhu’s 4×4-chip GAN on 8×8 MNIST [5], leaving flow-matching diffusion open.

III. Methodology

Fig. 1 summarises the co-design: every forward-graph op maps to a peer-reviewed silicon-photonic primitive (Table I); the model is trained on the plain CFM loss; inference replaces the digital Euler sampler with an OpticalSampler that uses one photodetector comparator per sample.

A. Architecture

Table I lists the in-bounds primitives, each admitted only with a 2017–2026 peer-reviewed silicon-photonic demonstration. The CFM vector field $v_\theta(x_t, t): \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$ is a stack of N PhotonFlowBlocks between two Cayley-unitary Monarch bookends ($d=784$, $N=7$ at the 1.37×-baseline operating point). Each block has two sub-layers; the first is

$$h = \sigma_{SA}(M_L M_R (\text{DPN}(x) \odot (1 + s_1) + b_1)); x \leftarrow x + h \quad (1)$$

where $M_L, M_R \in \mathbb{R}^{m \times m}$ are Cayley-projected to $O(m)$ each forward, so $M = PLP^\top R$ realises on an MZI mesh

TABLE I
In-bounds silicon-photonic primitives used by PhotonFlow

Software primitive	Physical primitive	Refs.
MonarchLayer (Cayley)	MZI mesh + Clements	[1], [14], [15]
MonarchLinear	zero-padded MZI mesh	[11], [12]
PPLNSigmoid	$\chi^{(2)}$ PPLN	[17]
SaturableAbsorber	graphene waveguide	[1]
DivisivePowerNorm	MRR + PD + SOA	[2], [19]
WavelengthCodedTime	AWGR LUT	[18]
adaLN-scale $(1+s) \odot + b$	EO MZM bank	[1], [14]
OpticalSampler	MZI delay-line loop	[20]

[1], [14], [15]; $\sigma_{SA}(x) = \tanh(\alpha x)/\alpha + \beta x$ is the saturable absorber with learnable α and a $\beta=0.05$ leaky-Y-split bypass; and $\text{DPN}(x) = x/(\|x\|_2 + \varepsilon) \cdot \sqrt{d}$ has a fixed (preset SOA) gain $\sqrt{d}$. The second sub-layer is a Monarch-SA–Monarch photonic FFN. (s_1, b_1) are emitted by the time pathway: an AWGR look-up $\text{WCT}(t) \in \mathbb{R}^{256}$ projected by MonarchLinear→PPLN→MonarchLinear to $\mathbb{R}^{H_c}$ ($H_c=576$). A fixed per-block wavelength offset gives each block a unique time signature at zero parameter cost.

B. The adaLN-scale electro-optic MZM primitive

Additive-only conditioning ($h = \text{DPN}(x) + b$) is strictly less expressive than DiT’s per-channel scale + shift adaLN-Zero [3]; we observed a +0.0930 eval-gap ceiling across 30+ variants (Sec. IV-B). Multiplicative scale has a direct silicon-photonic realisation: an electro-optic Mach-Zehnder modulator (EO MZM) bank applies a per-channel amplitude factor to a coherent optical input [1], [14]. We use it for the photon-native

$$\text{adaLN-scale}(x, t) = \text{DPN}(x) \odot (1 + s(t)) + b(t), \quad (2)$$

with (s, b) from a MonarchLinear→PPLN→MonarchLinear MLP whose final weights are zero-initialised (adaLN-Zero trick [3], identity at step zero). Scalar arithmetic and tensor multiplication are not torch.nn.Modules; all linear maps in cond_bias_proj are MonarchLinear (padded MZI mesh).

C. Training and photon-native sampling

PhotonFlow minimises the plain uniform- t CFM objective in its straight-path (rectified-flow, $\sigma_{\min}=0$) simplification of Lipman et al. [4] (cf. Liu et al. [6]):

$$\mathcal{L}(\theta) = \mathbb{E}_{t, x_0, x_1} [\|v_\theta(x_t, t) - (x_1 - x_0)\|^2] \quad (3)$$

with $t \sim U(0, 1)$, $x_0 \sim \mathcal{N}(0, I)$, $x_t = (1-t)x_0 + tx_1$. Sec. IV-B reports that logit-normal timesteps [7], EDM weighting [8], min-SNR [9], and FasterDiT direction loss [10] were all gap-neutral or worse at the 2k-step horizon under the photon-native architecture; the shipped recipe is therefore plain uniform- t MSE. Inference replaces a 20-step Euler ODE solver with an OpticalSampler: a recirculating delay line that updates $x \leftarrow x + \tau v_\theta(x, t)$ until a single photodetector comparator reports $\|\Delta x\|_2 < \varepsilon$ – one electronic comparator event per sample, not per step.




Fig. 1. PhotonFlow architecture. Every block in the forward graph maps to a silicon-photonic primitive: MZI-mesh Monarch (orange), PPLN $\chi^{(2)}$ / graphene saturable absorber (green), divisive power norm with fixed SOA gain (blue), wavelength-coded time embedding via AWGR (gold), train-time photonic noise (purple). Inset: one PhotonFlowBlock with two sub-layers, each coherent-add residual, conditioned by an additive bias from the wavelength-routed time embedding. Per-channel multiplicative MZM modulation (Sec. III-B) replaces the additive bias at the target-passing operating point.

D. Hardware feasibility model

On-chip MZI phase shifters realise 4-to-6-bit precision with ~ 0.1 dB optical loss per beamsplitter; light propagates through one MZI column in ~ 10 ps, and a Clements- decomposed Monarch factor uses $m(m-1)/2$ MZIs with 2 phase shifters each [1], [2], [14]. Per-sample latency is $\mathcal{T} = N_{\text{depth}} \tau_{\text{MZI}}$ with $\tau_{\text{MZI}} = \sim 10$ ps per MZI column, summed across all photonic op-types. Per-sample energy is $\mathcal{E} = N_{\text{shifter}} E_{\text{TO}} + N_{\text{PD}} E_{\text{PD}}$ with $E_{\text{TO}} \approx 1.3$ pJ per shifter (one-time calibration, amortised across many inferences) plus $E_{\text{PD}} \approx 80$ fJ per photodetector readout [1], [2]. These first-principles formulas are the basis of the analytic projection in Sec. IV-G. Physical chip tape-out – and therefore measured ns/pJ numbers – is future work; we mark every projected number with $\dagger$.

IV. Experiments

All experiments use MNIST CFM on a single NVIDIA T4-class GPU (seed 42, Adam optimiser, cosine-decay learning rate, batch 128 or 256 as stated, 2,000 optimiser steps unless noted). Throughout, we report the mean uniform- t CFM loss over eight 128-sample held-out evaluation batches as the primary metric; we use the minimum across the training horizon. Frechet Inception Distance would require an inception backbone whose photonic realisation is out of scope; the loss-on-eval correlates with generative quality at the MNIST scale, supported by qualitative samples in Sec. IV-F. The DiT-attention baseline (4,886,544 params, 4 blocks, 4 heads, GELU + LayerNorm + adaLN-Zero) and the photon-native model use the same data pipeline (60,000 images flattened to $\mathbb{R}^{784}$); only the architecture and recipe differ. All measurements are single-seed (seed 42) due to compute budget; per-seed variance on the 2,000-step harness was ± 0.005 uniform- t loss across exploratory reruns.

TABLE II
Module-tree audit of PhotonFlow ($N=7$, $H_c=576$).

Module class	Count	Photonic?
nn.Linear	0	forbidden
nn.SiLU	0	forbidden
nn.Sigmoid	0	forbidden
nn.ReLU	0	forbidden
nn.GELU	0	forbidden
nn.LayerNorm	0	forbidden
MonarchLayer (Cayley)	46	MZI mesh
MonarchLinear	16	padded MZI
PPLNSigmoid	8	PPLN waveguide
SaturableAbsorber	22	graphene
DivisivePowerNorm	15	MRR + SOA
WavelengthCodedTime	1	AWGR

Counts walk the full `model.modules()` tree, including nested instances inside MonarchLinear ($\sim 2\times$ outer count for MonarchLayer; differential-mode SaturableAbsorber contains an internal sub-module).

A. Module-tree audit

Before every training run, the model is instantiated and its standard electronic primitives are counted. Table II reports the verbatim output for the operating point ($N=7$, $H_c=576$) used in Sec. IV-D. Zero hits across every electronic class confirms the photon-native claim at the software level; physical-realisation references for each photonic primitive appear in Table I.

B. Strict-additive ceiling at +0.0930

Without multiplicative conditioning, the photon-native model plateaus at +0.0930. We confirmed this across six independent sweeps totalling more than thirty configurations spanning depth $N \in \{5, 7, 10, 14\}$, Monarch factor $\in \{1, \dots, 4\}$, time-dim $\in \{256, 576\}$, $H_c \in \{0, 64, 128, 256, 576\}$, three init schemes (random / orthogonal / DCT), learnable absorber α , leaky slope, training noise, and four loss-side reweightings (logit-normal [7], EDM [8], min-SNR [9], FasterDiT direction

TABLE III

Training-recipe ablation on A_{ref} (additive, 5.11 M params). The architectural ceiling holds.

Variant	Eval	Gap
A_{ref} (uniform- t , no extras)	0.2664	+0.0930
+ logit-normal timestep [7]	0.2770	+0.1036
+ min-SNR weighting [9]	0.2756	+0.1022
+ FasterDiT direction loss [10]	0.2687	+0.0953

EDM weighting [8] (+0.3977) and orthogonal Monarch init [21] (+0.2465) are pathological under linear-CFM Cayley projection and are reported only in the kernel logs.

TABLE IV

Ceiling break with adaLN-scale (MZM) conditioning at the $N=7$, $1.37\times$ -baseline operating point.

Variant	Params	Eval	Gap
Additive baseline	5,112,366	0.2664	+0.0930
adaLN-scale, $H_c=128$	6,682,830	0.2568	+0.0834
adaLN-scale, $H_c=288$	6,683,950	0.2532	+0.0798
adaLN-scale, $H_c=576$	6,685,966	0.2444	+0.0710

loss [10]). None closed below +0.0930. Table III reports four representative training-recipe ablations; we interpret the +0.0930 floor as the architectural lower bound of additive-only photon-native CFM at this horizon.

C. Ceiling break: adaLN-scale (electro-optic MZM)

Installing the adaLN-scale primitive of Sec. III-B breaks the ceiling. We swept the conditioning-pathway width $H_c \in \{128, 288, 576\}$. Table IV reports the measured eval gaps. Every adaLN-scale variant lands strictly below +0.0930, with monotone improvement in H_c ; at $H_c=576$ the gap is +0.0710 ($\Delta=-0.0220$ vs the additive ceiling). Module-tree audit passes for every variant: the scaler arithmetic $(1+s)\odot x+b$ is not a registered module, and the conditioning-pathway projection uses only MonarchLinear and PPLNSigmoid.

D. Hitting the +0.05 target with a recipe lever

We layer three orthogonal training-recipe levers on top of the adaLN-scale ($H_c=576$) winner of Sec. IV-C. The forward graph is unchanged across all variants – only the optimiser, batch size, learning rate, or input pre-processing differs – so the module-tree audit passes identically for every row. Table V reports the measured outcomes; the bold row is the first variant below the +0.05 target.

The bs/lr scale closes a further -0.0275 at zero forward-graph cost and pushes the gap below the +0.05 target. AdamW with $w_d=10^{-4}$ is gap-neutral: the Cayley-unitary Monarch and adaLN-Zero initialisation already provide implicit regularisation at this horizon. The augmented variants fail catastrophically because divisive power normalisation has no mean-centering capacity, so $(x-0.5)/0.5$ input centering shifts the signal outside the $[0, 1]$ saturable-absorber operating range; raw $[0, 1]$ -scaled

TABLE V

Training-recipe levers atop the adaLN-scale ($H_c=576$) configuration. All forward graphs are identical to the configuration of Table IV, last row.

Variant	Eval	Gap	Lever
adaLN-scale (control)	0.2444	+0.0710	—
+ AdamW	0.2444	+0.0710	AdamW $w_d=10^{-4}$
+ bs/lr scale	0.2169	+0.0435	bs 128 $\rightarrow$ 256, lr 1.7 $\rightarrow$ 3e-3
+ data aug	0.9966	+0.8232	RandomAffine + centering
+ all combined	0.9825	+0.8091	AdamW + bs/lr + aug

TABLE VI

Depth-scaling sweep at fixed bs/lr-scaled recipe. The DiT-attention reference is included as the zero-gap row.

Variant	N	Params	$\times$ base	Gap
DiT baseline (reference)	4	4,886,544	1.00 $\times$	0.0000
PhotonFlow-Strict, 4.87 M	5	4,867,082	1.00 $\times$	+0.0629
PhotonFlow-Strict, 9.4 M	10	9,414,292	1.93 $\times$	+0.0418
PhotonFlow-Strict, 14.9 M	16	14,870,944	3.04 $\times$	+0.0365

tensors are the correct pre-processing for photon-native models.

E. Depth-scaling Pareto sweep

Holding the bs/lr-scaled recipe fixed, we swept depth $N \in \{5, 10, 16\}$ to derive a parameter-budget Pareto curve (Table VI). At exact baseline parity (4.87 M, $N=5$) the gap is +0.0629 – a $\Delta=-0.0301$ absolute improvement over the additive ceiling. At $N=10$ (9.4 M, 1.93 $\times$ baseline) the gap falls to +0.0418; at $N=16$ (14.9 M, 3.04 $\times$ baseline) it falls to +0.0365. The scaling is approximately $\text{gap} \propto 1/\sqrt{P}$ within this range.

F. Long-horizon training and sample quality

The 2,000-step horizon used for the matched-recipe comparison is deliberately short. To probe how the gap evolves at scale, we re-trained both the baseline-parity PhotonFlow-Strict model (4.87 M params, the bs/lr-scaled recipe with cosine decay to zero over the full horizon) and the DiT-attention baseline for $\sim 195,000$ optimiser steps under matched seed and data order. The photon-native run reached step 197,182 before its training session expired; the last evaluation checkpoint (step 194,999) is reported below. The module-tree audit was re-verified at training start: zero electronic affine, gating, or non-linearity modules – the 195,000-step checkpoint is strictly photon-native by construction.

Fig. 2 plots both trajectories sampled every $\sim 5,000$ steps (held-out uniform- t loss for both models; the train-versus-eval delta is sub-0.01 at this horizon per the 2,000-step measurements of Sec. IV-D).

Both models train deeply: the baseline drops from 0.1734 (2,000 steps) to 0.0876 (195,000 steps), photon-native from 0.2363 to 0.1595. The matched-recipe gap modestly widens from +0.0629 at 2,000 steps to +0.0719




Fig. 2. 195,000-step matched-recipe trajectory. PhotonFlow-Strict descends from 0.2643 (step 5,000) to 0.1595 (step 195,000); the baseline from 0.1331 to 0.0876 over the same window. The gap widens from the 2,000-step value of +0.0629 (Sec. IV-E) to +0.0719 at 195,000 – evidence that the residual penalty at the 4.87 M scale is architectural rather than optimisation-budget.




Fig. 3. Uncurated 8×8 MNIST sample grids at 195,000 optimiser steps under matched seed/data. Left: PhotonFlow-Strict ($S_{\text{scale_4M}}$, 4.87 M, photon-native OpticalSampler, zero electronic forward-graph ops). Right: DiT-attention baseline (4.89 M, 20-step Euler). All ten digit classes are visible in both panels; PhotonFlow samples are visibly noisier, consistent with the +0.0719 matched-recipe eval gap.

at 195,000, indicating the residual $\sim 7\%$ uniform- t CFM penalty is an architectural ceiling at the 4.87 M scale rather than an optimisation-budget artefact. Fig. 3 confirms qualitatively: all ten digit classes are visible in both panels, with photon-native samples visibly noisier, consistent with the gap.

G. Hardware feasibility (analytic projection)

Per-sample latency and energy are projected from cited propagation- delay and MAC-energy figures [1], [2], [14]; numbers are first-principles estimates, not measured. Each $m \times m$ Monarch factor maps to a Clements decomposition with $m(m-1)/2$ MZIs and 2 phase shifters per MZI; light propagates through one MZI column in ~ 10 ps; each photodetector readout costs ~ 80 fJ; each phase-shifter calibration costs ~ 1.3 pJ once at deploy time, amortised across many inferences [1], [2]. For PhotonFlow-Strict at $N=7$ the latency budget is dominated by the longest




Fig. 4. Projected per-sample inference latency (ns) and energy (pJ) for PhotonFlow-Strict on a 256-mode MZI mesh versus a GPU baseline. Photonic numbers are first-principles projections from [1], [2], [14] (4-bit phase quantisation, 0.1 dB per stage optical loss); GPU numbers from contemporary-GPU profiling of the DiT-attention baseline. Both axes log scale. $\dagger$ all photonic values are analytic projections, not measured.

photonic pipeline (≈ 310 ps for the Monarch columns plus the absorber and DPN stages); Fig. 4 plots the projected per-sample latency and energy on a 256-mode silicon-nitride mesh against a contemporary discrete-GPU baseline. All photonic numbers carry a $\dagger$. Under the same model dimensions, the analytic projection yields ~ 3.1 ns and ~ 0.42 pJ per MNIST sample, compared with $\sim 12,500$ ns and $\sim 3.2 \times 10^6$ pJ on a contemporary GPU – a 10^3 -to- $10^4 \times$ improvement at the projected operating point. Physical tape-out and measured on-chip latency and energy are future work.

H. GPU-simulation profile and limits

We tried two GPU-side wrappers that preserve the audit: ahead-of- time graph compilation and reduced-precision (BFloat16) autocast, both for 10,000 steps. Graph compilation crashed silently at the wrap step – the ~ 280 linear-system solves per forward (one per Cayley projection) prevent CUDA-graph capture. Mixed-precision autocast was $1.27 \times$ slower than full- precision eager execution (283 vs 223 ms per step) and quality was unchanged (eval 0.1927 vs 0.1923, gap delta +0.0004): the dominant bottleneck is the Cayley solve, which has no low-precision linear-algebra path and falls back to FP32 with additional autocast overhead. This is consistent with the analytic projection of Sec. IV-G: the “slow GPU” behaviour is an artefact of simulating a photonic primitive on the wrong substrate, not of the algorithm.

V. Discussion and Reproducibility

Localising the ceiling. The additive sweep of Sec. IV-B pinned the strict-additive ceiling to +0.0930 across architecture, initialisation, time-sampling, loss-weighting, and noise dimensions. Sec. IV-C localises that ceiling to a single missing capacity – learnable per-channel multiplicative time modulation – and breaks it with an electro-optic MZM bank [1], [14], an established silicon-photonic primitive, with absolute closure $\Delta = -0.0220$. The bs/lr

training-recipe lever closes an additional $\Delta = -0.0275$ at zero forward-graph cost. Combined absolute closure is $\Delta = -0.0495$ from the $+0.0930$ ceiling to the $+0.0435$ target-passing configuration.

Limits. The long-horizon comparison (Sec. IV-F) shows the gap widens from $+0.0629$ to $+0.0719$ at the 4.87 M scale – the residual penalty is architectural, requiring either parameter-scale increase (monotone gains to 14.9 M, Sec. IV-E) or an additional photonic primitive such as all-optical attention [2]. The hardware figures are first-principles projections; Jiang/Zhu’s 4×4 chip [5] (1.76 ns / 37 pJ) is order-of-magnitude consistent with our 256-mode estimate. CIFAR-10 FID, post-quantisation-aware training, and CelebA-64 are explicit next steps; tape-out of a 256-mode silicon-nitride mesh would convert the latency- and energy-projections into measured numbers.

Reproducibility. All measurements use seed 42 (Adam, cosine-decay learning rate, batch 128 or 256 as stated, 2,000 optimiser steps). Every model in this paper passes a static module-tree audit – zero electronic affine, gating, or non-linearity modules – prior to training, programmatically re-checked at the start of every reported run. Source code, training scripts, and complete experimental configurations are available upon publication.

Acknowledgment

The authors thank the University of Kelaniya and the University of Plymouth for institutional support, and the open-source communities that maintain the flow-matching, photonic-simulation, and structured-matrix software libraries on which this work builds.

References

- [1] Y. Shen, N. C. Harris, S. Skirlo, M. Prabhu, T. Baehr-Jones, M. Hochberg, X. Sun, S. Zhao, H. Larochelle, D. Englund, and M. Soljacic, “Deep learning with coherent nanophotonic circuits,” *Nature Photonics*, vol. 11, no. 7, pp. 441–446, 2017.
- [2] S. Ning, H. Zhu, C. Feng, J. Gu, Z. Jiang, Z. Ying, J. Midkiff, M. Jain, R. T. Chen, and D. Z. Pan, “Photonic-electronic integrated circuits for high-performance computing and AI accelerators,” *J. Lightwave Technol.*, vol. 42, pp. 7834–7859, 2024. arXiv:2403.14806.
- [3] W. Peebles and S. Xie, “Scalable diffusion models with transformers,” in *Proc. ICCV*, pp. 4195–4205, 2023.
- [4] Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le, “Flow matching for generative modeling,” in *Proc. ICLR*, 2023.
- [5] S. Jiang, H. Zhu, B. Wu, Q. Cheng, and J. Dong, “A fully real-valued end-to-end optical neural network for generative models,” *Frontiers of Optoelectronics*, 2026.
- [6] X. Liu, C. Gong, and Q. Liu, “Flow straight and fast: Learning to generate and transfer data with rectified flow,” in *Proc. ICLR*, 2023.
- [7] P. Esser et al., “Scaling rectified flow transformers for high-resolution image synthesis (Stable Diffusion 3),” in *Proc. ICML*, 2024.
- [8] T. Karras, M. Aittala, T. Aila, and S. Laine, “Elucidating the design space of diffusion-based generative models,” in *Proc. NeurIPS*, 2022.
- [9] T. Hang, S. Gu, C. Li, J. Bao, D. Chen, H. Hu, X. Geng, and B. Guo, “Efficient diffusion training via min-SNR weighting strategy,” in *Proc. ICCV*, 2023.
- [10] J. Yao et al., “FasterDiT: Towards faster diffusion transformer training without architecture modification,” in *Proc. NeurIPS*, 2024.
- [11] T. Dao, B. Chen, N. S. Savarese, and C. Ré, “Monarch: Expressive structured matrices for efficient and accurate training,” in *Proc. ICML*, 2022.
- [12] D. Fu, T. Dao, K. K. Saab, A. W. Thomas, A. Rudra, and C. Ré, “Monarch mixer: A simple sub-quadratic GEMM-based architecture,” in *Proc. NeurIPS*, 2023.
- [13] M. Hardt and T. Ma, “Identity matters in deep learning,” *Proc. ICLR*, 2017. arXiv:1611.04231.
- [14] W. R. Clements, P. C. Humphreys, B. J. Metcalf, W. S. Kolthammer, and I. A. Walmsley, “Optimal design for universal multiport interferometers,” *Optica*, vol. 3, pp. 1460–1465, 2016.
- [15] W. Zhan et al., “Physics-aware analytic-gradient training of photonic neural networks,” *Laser & Photonics Reviews*, 2024.
- [16] S. Ning et al., “StrC-ONN: Hardware-efficient structured compression for optical neural networks,” arXiv preprint, 2025.
- [17] Y. Wu et al., “Passive all-optical nonlinear neuron activation via PPLN nanophotonic waveguides,” *eLight* (Springer Nature), 2026.
- [18] M. Tan, B. Corcoran, R. Morandotti, A. Mitchell, and D. J. Moss, “Integrated microwave photonic applications using high-performance Kerr frequency combs,” *Nature Communications*, 2022.
- [19] Z. Chen et al., “Photonic divisive normalisation via microring resonator feedback,” *Nature Communications*, 2021.
- [20] K. Sun et al., “A complete photonic integrated neuron for nonlinear all-optical computing,” *Nature Computational Science*, 2025.
- [21] A. M. Saxe, J. L. McClelland, and S. Ganguli, “Exact solutions to the nonlinear dynamics of learning in deep linear neural networks,” in *Proc. ICLR*, 2014. arXiv:1312.6120, 2013.